

Parallel Huffman Coding

Final Presentation

Chen Luo

Patrick (Zhanxiang) Huang

Background

- What is Huffman Coding?
 - Compression algorithm to generate optimal prefix code.
- Why parallelize Huffman Coding is important?
 - Sequential Huffman Coding is used by many compression libraries
 - Huffman Coding takes a significant percentage of compression time

Sequential Compression

- Step 1: Build Symbol Frequencies Histogram

Input File:

0x01	0x02	0x00	0x01
------	------	------	------




Symbol	Freq
0x00	1
0x01	2
0x02	1

Sequential Compression

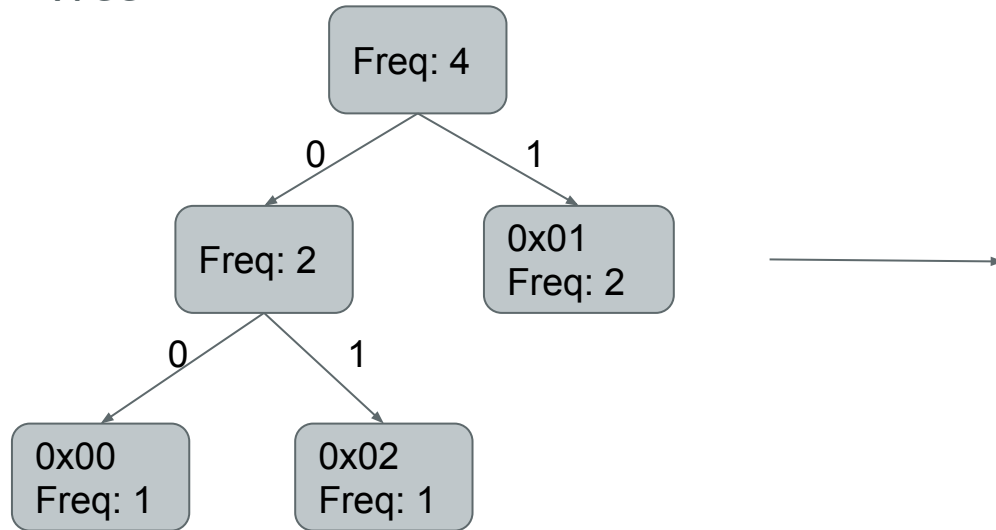
- Step 2: Build Huffman Tree

Symbol	Freq
0x00	1
0x01	2
0x02	1



Sequential Compression

- Step 3: Build Prefix Code Table from Huffman Tree



Byte	Code
0x00	00
0x01	1
0x02	01

Sequential Compression

- Step 4: Encode file using Prefix Code

Input File

0x01	0x02	0x00	0x01
------	------	------	------

Compression →

Output File

header	...	header	101001
--------	-----	--------	--------


Byte	Code
0x00	00
0x01	1
0x02	01

Compression Ratio : $6 / 32 = 0.1875$

Serialized Symbol table	Chunk Start Offset	Compressed Bits
----------------------------	-----------------------	--------------------


Sequential Decompression

- Rebuild Huffman Tree




Sequential Decompression

- Decompress file using Huffman Tree



Sequential Compression Bottlenecks

- 78% of the Time Spends on the Second Pass of the Data to Do Compression
- 16% of the Time Spends on Generating Symbol Histogram
- 3% of the Time Spends on Building Huffman Tree



Sequential Decompression Bottlenecks


- 0.2% of the time is spent on building Huffman Tree
- 99.8% of the time is spent on decoding files

Our Solution

- Compression
 - Build Symbol Frequencies Histogram (Parallel)
 - Build Huffman Tree and Build Prefix Code Table (Sequential)
 - Compress File using Prefix Code Table (Parallel)
- Decompression
 - Build Huffman Tree (Sequential)
 - Decompress File by traversing Huffman Tree (Parallel)

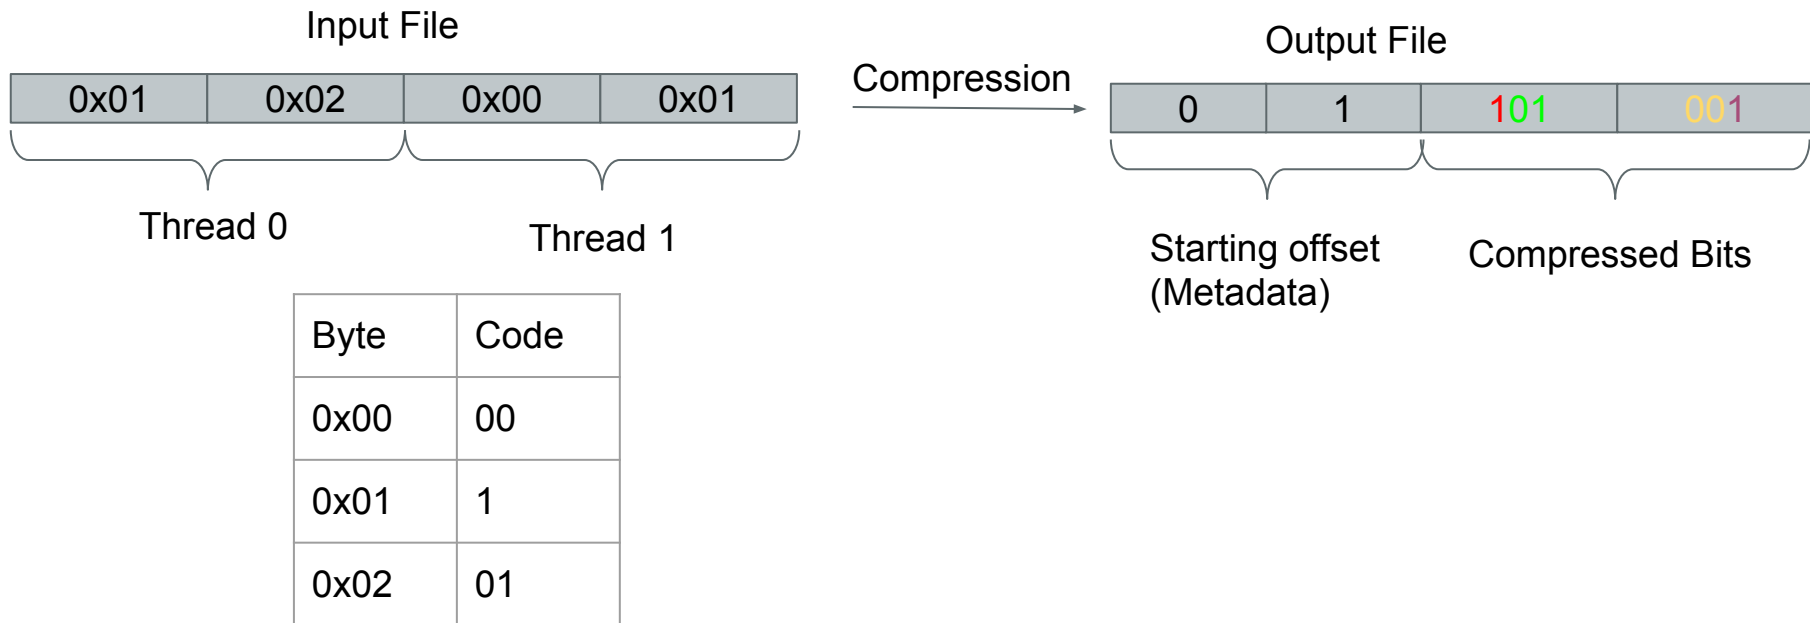
Parallel Compression

- Step 1: Build Symbol Frequencies Histogram



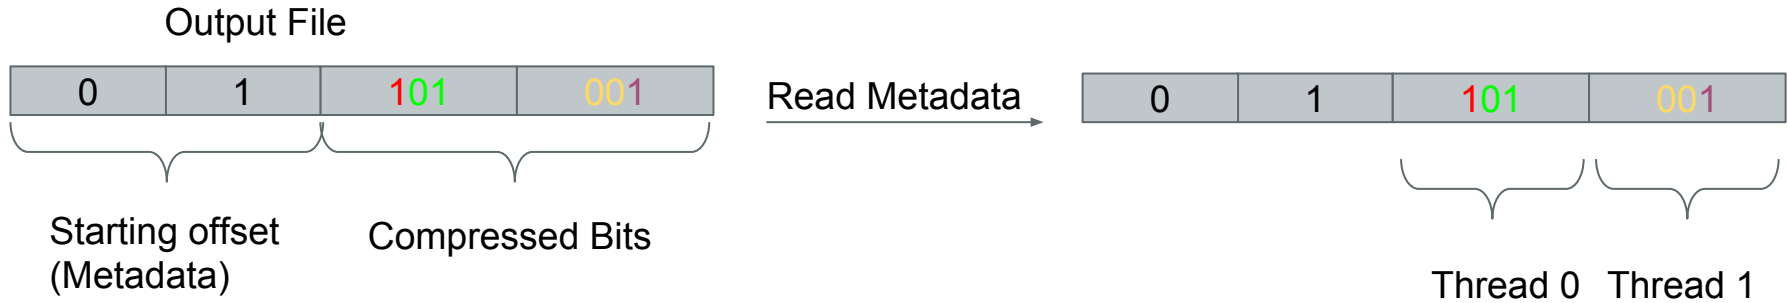
Parallel Compression

- Step 4: Encode file using Prefix Code



Parallel Decompression


- Step 2: Decompress File using Huffman Tree



Evaluation Setup

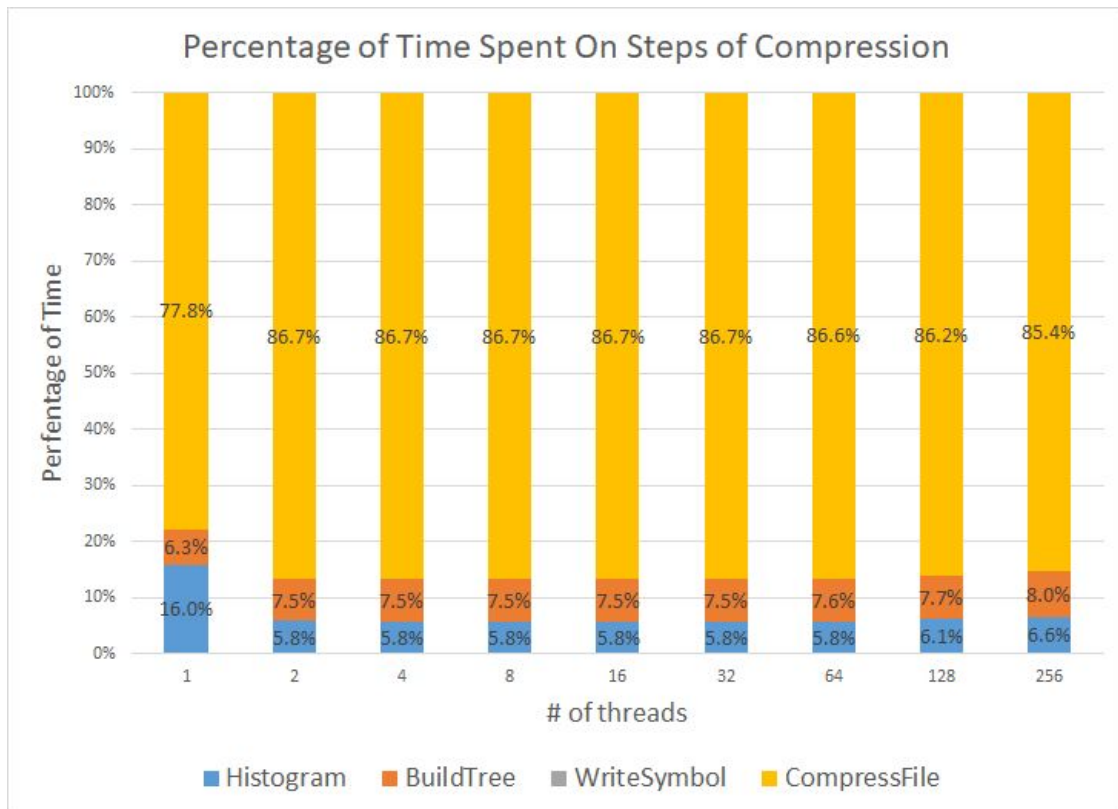
- Xeon Phi (KNL) Co-processor
 - 68 Cores, 256 threads
- Xeon E5-2699 v4 @2.20GHz
 - 88 Cores, NUMA Architecture (22 Cores per socket)

Parallel Encoding Only





Xeon Phi (KNL) Co-processor (68 Cores, 256 Threads), 5.5GB Wiki dataset

Parallel Encoding + Parallel Histogram



Xeon Phi (KNL) Co-processor (68 Cores, 256 Threads), 5.5GB Wiki dataset

Speedup on Xeon Phi – Linear Speedup

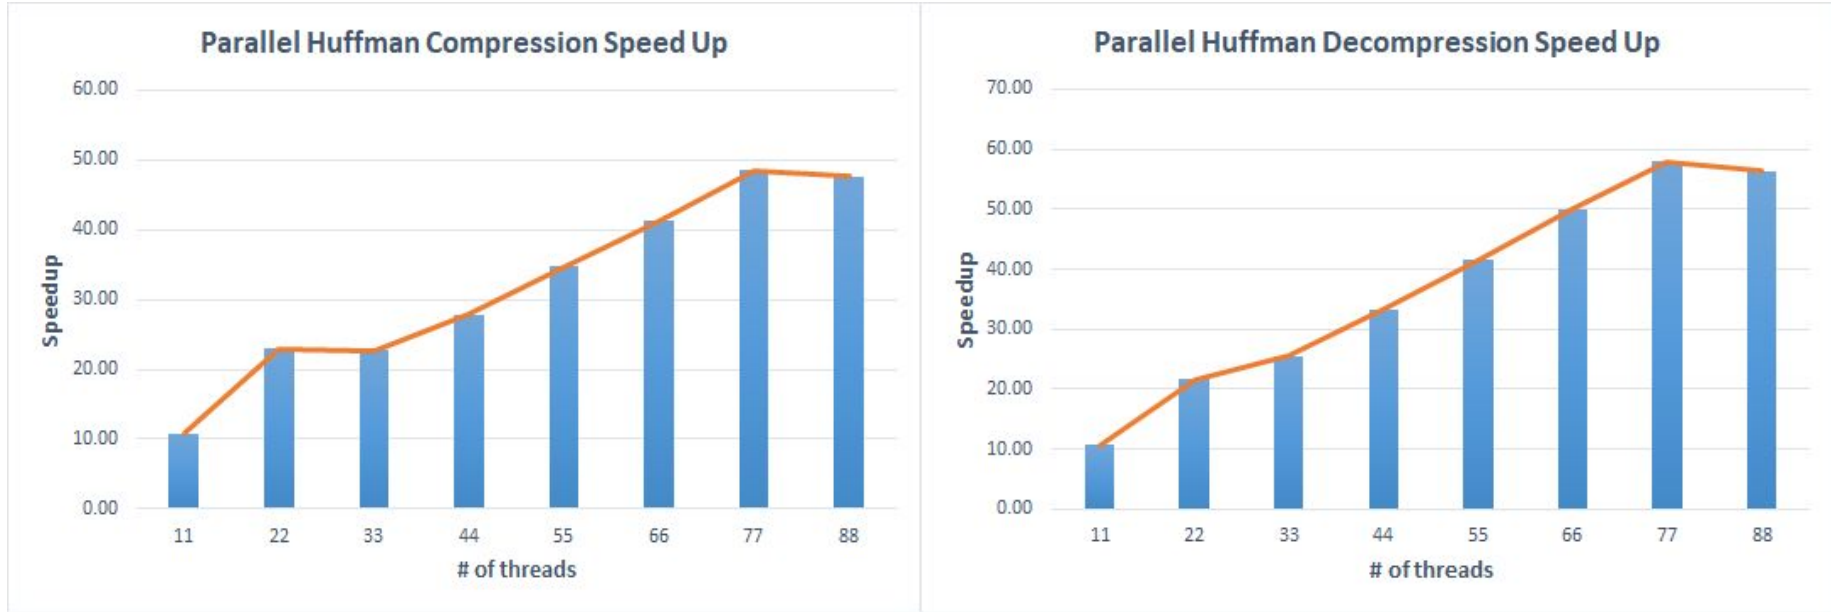


Xeon Phi (KNL) Co-processor (68 cores, 256 threads)

Speedup Analysis


- Memory Bandwidth
 - Symbol list is $256 * 24 = 6144$ bytes. Reading sequentially from input file
 - Total Working Set Size: ~6kb. L1 Cache Size: 64 kb
- Workload Balance
 - Each thread finishes roughly at the same time. Only 5% difference.

Speedup on Xeon E5-2699 (NUMA)



Xeon E5-2699 v4 @2.20GHz (88 Cores, 4 NUMA Socket)

Speedup on Xeon E5-2699 (NUMA)



Xeon E5-2699 v4 @2.20GHz (88 Cores, 4 NUMA Socket)

Future Work

- Modify our compression algorithm to work better on NUMA architecture
 - Each thread reads a file chunk into its local memory.
- Adapt similar techniques to other compression algorithms


Other Alternatives

- ISPC and SIMD instructions
 - Huffman Compression is not a perfect workload for SIMD
 - Why?

Bit-level Conflicts

- ISPC and SIMD instructions
 - Huffman Compression is not a perfect workload for SIMD
 - Why?

Number of Bits




Bytes Offset

1 2 3 4

SIMD Unit

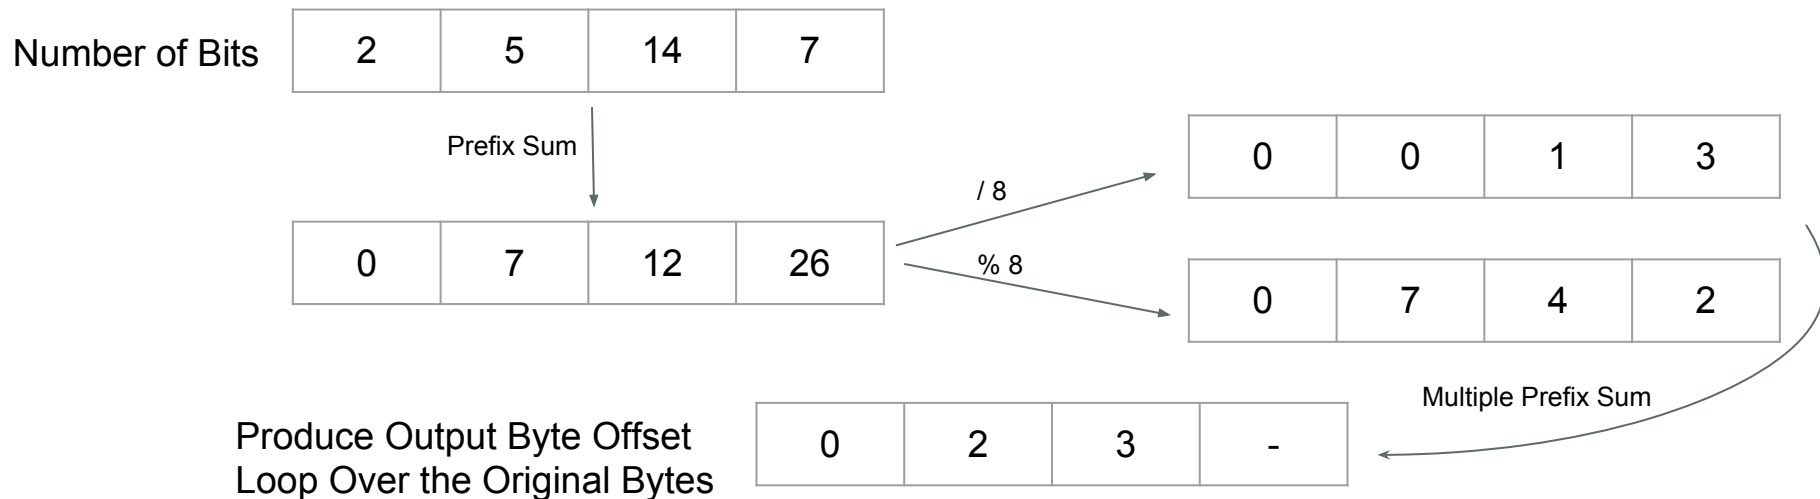


Output Bytes





More instructions

- ISPC and SIMD instructions
 - Handle structs is tricky
 - Want each SIMD write independent bytes, but require many instructions



Backup Slide: Speedup on Xeon Phi



Xeon Phi (KNL) Co-processor (68 cores, 256 threads)